

# Project 2 Documentation

John Carlo Sumabat

The following program is a simple threaded implementation of a file server in C. The implementation level is Level 4. The entirety of the program is enclosed in `file_server.c`. You may access the **documentation video** by opening [this link](#).

## 1 Introduction

### 1.1 Implementation Overview

This section provides a summary of the mechanism with which the program works. Future sections will elaborate on the aspects discussed here.

The system mainly relies on two arrays to work: one that keeps track of “active” files (named `files`), and a corresponding array that keeps data on each entry in `files` (named `jobs`). Concurrency is achieved via a thread pool (an array named `tid`) and a corresponding array that keeps track of which threads are active (named `threads`).

At any point during the program’s execution, at most `N` files may receive operations and at most `M` threads may be active at one time, where `N` and `M` are constants that can be changed prior to compilation. These two constants reflect the sizes of `files` and `tid`, respectively.

User input is parsed by the master thread. Upon receiving a command, it performs various checks to determine the correct data to provide an available worker thread from `tid[]` (elaborated in [Section 2.7](#)). This data is stored in a `struct Jobdata` whose pointer is passed as a parameter to `pthread_create()` which dispatches an available worker thread from the thread pool. The worker thread performs its assigned operation, allowing the master thread to immediately receive input again. Since the execution of an operation is done on a different thread, the master thread can continuously receive requests.

One such check that the master thread performs is identifying whether the file specified by the user-supplied filepath is “active”. A file is deemed *active* if at least one worker thread will perform a read or write operation on it. If a file is active, then its filepath is stored in `files` and it has its own lock independent of other files. An active file has each of its worker threads assigned a job ID. A thread’s job ID is determined by their order of arrival, such that they work on a file in a “First In, First Out” (FIFO) manner. That is, a thread can only do its assigned operation if its job ID is equal to the current job that the file is accepting. This makes it so that sequential reads or writes to the same file happen in the order that they are requested.

Such enforcement of order is made possible by a file’s corresponding `struct TJob` which records an active file’s activity (elaborated in [Sections 2.2](#) and [2.6](#)). The pointer to this struct is located in the `jobs` array, at the same index where a file’s pointer is located in the `files` array.

Whereas multiple requests to the same file are executed sequentially, multiple requests to different files are handled concurrently. This is because any two files with different filepaths are tracked via different indices in the `files` array. Since each active file has its own lock, locking

one file does not interrupt operations on other files. As such, operations on different files happen simultaneously (but no order is enforced).

Deadlocks and race conditions are obstacles that are present with this implementation. Details on how they are prevented are found throughout future sections.

## 1.2 Rationale

The goal of this implementation is to allow concurrent operations on multiple files. For ease of development, dynamic management of jobs via nodes and linked lists was not pursued. Rather, a static approach was taken via arrays whose size can be tweaked based on anticipated workload. By allowing each active file to occupy a single space in `files` (and therefore `jobs`), the problem boils down to properly managing locks and job logic.

# 2 Implementation Details

## 2.1 Macros

Three constants are defined (Code Block 1). `N` and `M` are the maximum sizes of the `files` and `tid` arrays, respectively, while `MAXLEN` stands for the maximum length of any user input.

---

**Code Block 1** User-defined constants

---

```
8 #define N 500           // max number of files that can be used at a time
9 #define M 1000          // max number of threads that can run at a time
10 #define MAXLEN 120      // max length of any file server command
```

---

## 2.2 Structs

Two structs are defined. The first one is `struct TJob` (Code Block 2). As mentioned in Section 1.1, this holds data about each active file. The `workers` member records the number of threads pending execution. Once this value becomes zero, the `files` and corresponding `jobs` entries can be safely cleared. The `count` member increments every time a new worker thread wishes to work on a file. The resulting value is assigned to the worker thread's `struct Jobdata` as its `job_ID`. The `curr_job` member indicates the `job_ID` of the worker thread whose operation it is currently accepting.

The struct also has a condition variable and two mutexes. The condition variable `done` is broadcast every time a worker thread finishes an operation to wake up other waiting threads. The mutex `wlock` is used when updating the `worker` count, while the mutex `global_lock` is used whenever a worker thread actually performs its operation. Both of these mutexes are required to prevent race conditions since more than one thread may want to interact with a file's `struct TJob` at the same time.

The second struct is `struct Jobdata` (Code Block 3). This holds the data that each worker thread uses to work on a job, namely:

- `cmdidx`: the assigned command (i.e., read, write, or empty)
- `fileidx`: the index of the assigned file in the `files` array
- `job_ID`: its position in a file's "queue" of workers
- `worker_ID`: its index in the `tid` array
- `string`: information to write to the file, if any

---

#### Code Block 2 struct TJob definition

---

```
17 typedef struct TJob {
18     int workers;
19     int count;
20     int curr_job;
21     pthread_cond_t done;          // condition variable for finished jobs
22     pthread_mutex_t wlock;        // lock for updating worker count
23     pthread_mutex_t global_lock; // lock for reading/writing
24 } tjob;
```

---

---

#### Code Block 3 Struct Jobdata definition

---

```
30 typedef struct Jobdata {
31     int cmdidx;
32     int fileidx;
33     int job_ID;
34     int worker_ID;
35     char string[51];
36 } jobdata;
```

---

### 2.3 Global Variables

**Code Block 4** shows the global variables used in the program. Note that the first four arrays are the arrays mentioned in **Section 1.1**. The three file pointers are used for the log files `commands.txt`, `read.txt`, and `empty.txt`.

The mutex array `tlock` provides a lock for every entry in `files` and `jobs`, while `qlock` provides a single lock for the thread pool. One mutex is required per entry in `files` to enable concurrent operations (since having only one lock would cause blocking, defeating the purpose of threads). On the other hand, the act of dispatching and exiting threads does not need to be concurrent (only thread-safe) so one lock is sufficient. Lastly, `rlock` and `elock` are mutexes for `read.txt` and `empty.txt`, respectively. Since more than one thread may want to write to the said log files at the same time, mutexes are necessary to prevent write errors.

Next, the `cmds` array is used to help check the validity of user input. The `space` array is used as a parameter to the `strtok_r()` function to tokenize user input. Lastly, `testmode` is a flag that allows the program to print test statements when set to `1`.

### 2.4 Test Functions

In order to verify the correctness of the program, several test functions were created.

- The function `pfiles()` (**Code Block 5**) prints all active files. An entry in `files` has an active file if the content at an index `i` is not " ".
- The function `pft()` (**Code Block 6**) prints all active worker threads. An entry in `threads` has an active worker if the content at an index `i` equals `1`.
- Every time a worker thread is dispatched, `pwork()` (**Code Block 7**) may be called to print the members from its corresponding `struct Jobdata`.
- Lastly, `pjob()` (**Code Block 8**) prints the members of a file's corresponding `struct TJob` given the file's index in the `files` array.

These functions are called via `test_print()` (**Code Block 9**) after every user request if the

---

#### Code Block 4 Global variables

---

```
40 char files[N][MAXLEN];           // keeps track of currently active jobs
41 tjob *jobs[N] = { 0 };           // contains data of jobs
42 pthread_t tid[M];                // worker threads
43 int threads[M] = { 0 };          // keeps track of currently working threads
44 FILE *cmdfile;                   // log file for commands
45 FILE *readfile;                  // log file for reads
46 FILE *emptyfile;                 // log file for empty
47 pthread_mutex_t tlock[N];        // lock for accessing jobs
48 pthread_mutex_t qlock;           // lock for accessing working threads
49 pthread_mutex_t rlock;           // lock for read.txt
50 pthread_mutex_t elock;           // lock for empty.txt
51
52 char cmds[][MAXLEN] = { "read", "write", "empty", };
53 char space[2] = " ";
54 int test_mode;
```

---

global variable `test_mode` equals 1. It takes in the index of a file in the `files` array and a thread's struct `Jobdata` as parameters.

---

#### Code Block 5 pfiles function

---

```
61 void pfiles() {
62     for ( int i = 0; i < N; i++ ) {
63         if ( strcmp( files[i], " " ) != 0 ) {
64             printf( "files[%d]: %s\n", i, files[i] );
65         }
66     }
67 }
```

---

---

#### Code Block 6 pft function

---

```
73 void pft() {
74     printf( "active threads: " );
75     for( int i = 0; i < M; i++ ) {
76         if( threads[i] == 1 ) {
77             printf( "%d, ", i );
78         }
79     }
80     printf( "\n" );
81 }
```

---

---

**Code Block 7 pwork function**

---

```
89 void pwork( jobdata *data ) {
90     printf( "thread working on fileidx = %d with:\n"
91         "\tcommand = %s\n"
92         "\tjob_ID = %d\n"
93         "\tworker_ID = %d\n"
94         "\tstring = %s\n"
95         ,
96         data->fileidx,
97         cmds[data->cmdidx],
98         data->job_ID,
99         data->worker_ID,
100         data->string );
101 }
```

---

---

**Code Block 8 pjob function**

---

```
109 void pjob( int idx ) {
110     tjob *job = jobs[idx];
111     printf( "file <<%s>> currently has:\n"
112         "\tworkers = %d\n"
113         "\tcount = %d\n"
114         "\tcurr_job = %d\n"
115         ,
116         files[idx],
117         job->workers,
118         job->count,
119         job->curr_job );
120 }
```

---

---

**Code Block 9 test\_print function**

---

```
126 void test_print( int fileidx, jobdata *data ) {
127     pfiles();
128     pft();
129     pjob( fileidx );
130     pwork( data );
131 }
```

---

## 2.5 Helper Functions

To help with repetitive tasks, some helper functions were created.

- The function `getcmd()` (Code Block 10) accepts a buffer and a buffer size. It sets the buffer to zero and places user input onto it.
- Every time a valid command is entered, `log_command()` (Code Block 11) takes the current time as a string, concatenates it with the user request, and appends it to `commands.txt`.
- The function `strindex()` (Code Block 12) searches the index of a string in a string array (the size of the array is also a parameter). It returns `-1` if the string is not found.

---

### Code Block 10 `getcmd` function

---

```
142 int getcmd( char *buf, int nbuf ) {
143     printf( "> " );
144     memset( buf, 0, nbuf );
145     if( fgets( buf, nbuf, stdin ) != NULL ) {
146         return 0;
147     }
148     return 1;
149 }
```

---

---

### Code Block 11 `log_command` function

---

```
157 void log_command( char *cmd ) {
158     time_t curr_time;
159     char log[150] = { 0 };
160     curr_time = time( NULL );
161
162     // add colon and space after timestamp
163     strcpy( log, ctime( &curr_time ) );
164     log[strlen( log ) - 1] = ':'; // turn '\n' to ':'
165     strcat( log, " " );
166
167     cmdfile = fopen( "commands.txt", "a" );
168     fputs( strcat( log, cmd ), cmdfile );
169     fclose( cmdfile );
170 }
```

---

---

### Code Block 12 `strindex` function

---

```
182 int strindex( char *str, char arr[][MAXLEN], int narr ) {
183     for( int i = 0; i < narr; i++ ) {
184         if( strcmp( str, arr[i] ) == 0 ) {
185             return i;
186         }
187     }
188     return -1;
189 }
```

---

## 2.6 Worker Thread Functions

When worker threads are dispatched, they call a few functions to ensure that proper order is enforced and race conditions are prevented. Recall that each thread has its own `struct Jobdata`, and a file that a thread is assigned to has a corresponding `struct TJob` for managing worker threads.

- The function `dispatch_worker()` (Code Block 14) is where all worker threads begin. Using the provided `struct Jobdata`, a thread calls `enqueue_worker()` to add itself to the `workers` member of the assigned file's `struct TJob`. This allows the file to stay active as long as at least one thread is still pending execution.

Afterwards, the worker thread is put to sleep via `xsleep()`. The `global_lock` of the file's `struct TJob` is then acquired to prevent concurrent operations on the same file. While the thread's job ID is not equal to the `struct TJob`'s current job, the thread goes to sleep.

When it is the worker thread's turn to operate on a file, its `struct Jobdata`'s `cmdidx` is checked via a `switch` statement to identify which operation it will execute (and thus which function to call). After doing its operation, the worker thread releases `global_lock`, removes itself from `worker` via `dequeue_worker()`, and calls `finish_job()`.

- Before a thread carries out an operation, `enqueue_worker()` (Code Block 15) is called to increment the `workers` member of the assigned file's `struct TJob`. Since multiple threads may want to call this function at one time, the struct's `wlock` is used to prevent race conditions.
- After an operation, `dequeue_worker()` (Code Block 16) is called to reverse the effect of `enqueue_worker()`.
- The function `xsleep()` (Code Block 13) generates a random number between 0–99. If the generated number is less than 80, the worker thread sleeps for 1 second. Otherwise, the sleep duration is 6 seconds. This way, the chances of a thread to sleep 1 and 6 seconds is 80% and 20%, respectively. If testing is enabled, the sleep duration is printed.

When performing an assigned operation, a worker thread calls one of three functions: `xread()`, `xwrite()`, or `xempty()` (corresponding to “read”, “write”, and “empty” commands). Each of these functions accepts a `struct Jobdata` which allows the thread to identify which file to open and what string to append, if any. None of these functions acquire or release `global_lock` when reading from or writing to the assigned file since the calling function `dispatch_worker()` does this beforehand.

- The function `xread()` (Code Block 17) writes the contents of a file to `read.txt`. It first acquires the global mutex `rlock` to prevent write errors, since multiple threads may call `xread()` at the same time. The function first tries to open the thread's assigned file. If the assigned file does not exist, it writes this fact in `read.txt`. Otherwise, the contents of the assigned file are copied to `read.txt`, character by character. Afterwards, both the assigned file and `read.txt` are closed and `rlock` is released. If testing is enabled, then the name of the read file is printed before returning.
- The function `xwrite()` (Code Block 18) writes a user-specified string to the thread's assigned file. First, it sleeps for  $25 \times n$  milliseconds, where  $n$  is the size of the string. It then attempts to open the assigned file. If the file does not exist, it will create it. However, if the directory of the file does not exist, it will *not* proceed. Otherwise, the string is appended to the file all at once via `fputs()`. Again, there is no need to lock the file within the function. If testing is enabled, then the name of the file and the sleep duration are printed before returning.

- The function `xempty()` (Code Block 19) writes the contents of a file to `empty.txt` then clears its contents. Similar to `xread()`, the global mutex `elock` is acquired to prevent write errors from simultaneous writes. If the assigned file does not exist or is already empty, it writes this fact in `empty.txt`. Otherwise, the contents of the assigned file are copied to `empty.txt`, character by character, then it is emptied via an `fopen()` call. Afterwards (or if the file exists but was found already empty), the file is closed and `elock` is released. It then goes to sleep for a random amount of time between 7 and 10 seconds via a `rand()` call. Since  $(n \bmod 4) \in \{0, \dots, 3\}$  for any integer  $n$ , the expression  $(r \bmod 4) + 7$  will provide a random number between 7 and 10 for a random integer  $r$ . This number is then used as a parameter to `sleep()` to make the thread sleep. If testing is enabled, then the name of the file and the sleep duration are printed before returning.

The last function that a worker thread calls is `finish_job()` (Code Block 20) which takes in the thread's `struct Jobdata` as a parameter. Since multiple threads may complete a job at the same time, the lock `tlock[i]` is acquired, where `i` is the index of the file in the `files` array.

The function first checks if there are worker threads pending execution by checking the `struct TJob`'s `workers` member. If it is equal to zero, then the `struct TJob` is freed, its entry in `jobs` is set to `NULL`, and the corresponding entry in `files` is set to `" "`. Otherwise, the `curr_job` member is incremented and the remaining worker threads are woken up via `pthread_cond_broadcast()`. Afterwards, `tlock[i]` is released.

In any case, the specific worker thread that called `finish_job()` will retire at this point. In order to prevent race conditions (e.g., the same thread retiring and being dispatched at the same time), `qlock` is acquired. The thread's entry in `threads` is set to `0`, and its corresponding `struct Jobdata` is freed.

---

#### Code Block 13 xsleep function

---

```

252 void xsleep( int worker_ID ) {
253     int s;
254
255     srand(time(0));
256     int r = rand()%100;
257     s = ( r < 80 ) ? 1 : 6;
258     sleep(s);
259
260     if( test_mode ) {
261         printf( "thread %d slept for %d seconds.\n", worker_ID, s );
262     }
263 }
```

---



---

#### Code Block 14 dispatch\_worker function

---

```
391 void *dispatch_worker( void *arg ) {
392     jobdata *data = (jobdata *) arg;
393     tjob *job = jobs[data->fileidx];
394
395     // begin operation
396     enqueue_worker( job );
397     xsleep( data->worker_ID );
398
399     pthread_mutex_lock( &(amp; job->global_lock ) );
400
401     while( job->curr_job != data->job_ID ) {
402         pthread_cond_wait( &(amp; job->done ), &(amp; job->global_lock ) );
403     }
404
405     // execute command
406     switch( data->cmdidx ) {
407         case 0:
408             xread( data );
409             break;
410         case 1:
411             xwrite( data );
412             break;
413         case 2:
414             xempty( data );
415             break;
416     }
417
418     pthread_mutex_unlock( &(amp; job->global_lock ) );
419
420     // end operation
421     dequeue_worker( job );
422     finish_job( data );
423 }
```

---

---

#### Code Block 15 enqueue\_worker function

---

```
230 void enqueue_worker( tjob *job ) {
231     pthread_mutex_lock( &(amp; job->wlock ) );
232     job->workers++;
233     pthread_mutex_unlock( &(amp; job->wlock ) );
234 }
```

---

---

#### Code Block 16 dequeue\_worker function

---

```
242 void dequeue_worker( tjob *job ) {
243     pthread_mutex_lock( &(amp; job->wlock ) );
244     job->workers--;
245     pthread_mutex_unlock( &(amp; job->wlock ) );
246 }
```

---

---

**Code Block 17** xread function

---

```
271 void xread( jobdata *data ) {
272     FILE *file;
273     char *fname;
274
275     pthread_mutex_lock( &rlock );
276
277     fname = files[data->fileidx];
278     readfile = fopen( "read.txt", "a" );
279     file = fopen( fname, "r" );
280
281     fprintf( readfile, "read %s: ", fname );
282
283     if( file == NULL ) {
284         fputs( "FILE DNE\n", readfile );
285     }
286     else {
287         char ch;
288         while(( ch = fgetc( file )) != EOF ) {
289             fputc( ch, readfile );
290         }
291         fputc( '\n', readfile );
292         fclose( file );
293     }
294
295     fclose( readfile );
296     pthread_mutex_unlock( &rlock );
297
298     if( test_mode ) {
299         printf( "read %s complete.\n", fname);
300     }
301 }
```

---

---

**Code Block 18** xwrite function

---

```
309 void xwrite( jobdata *data ) {
310     FILE *file;
311     char *fname;
312     char *string = data->string;
313     int ms = 25 * strlen( string );
314
315     usleep( ( unsigned int )( 1000 * ms ));
316
317     fname = files[data->fileidx];
318     file = fopen( fname, "a" );
319
320     if( file == NULL ) {
321         printf( "Target directory %s does not exist.\n> ", fname );
322     }
323     else {
324         fputs( string, file );
325         fclose( file );
326     }
327
328     if( test_mode ) {
329         printf( "write %s complete. slept for %d milliseconds.\n", fname, ms);
330     }
331 }
```

---

---

**Code Block 19** xempty function

---

```
339 void xempty( jobdata *data ) {
340     FILE *file;
341     char *fname;
342     char ch;
343
344     pthread_mutex_lock( &elock );
345
346     fname = files[data->fileidx];
347     emptyfile = fopen( "empty.txt", "a" );
348     file = fopen( fname, "r" );
349
350     fprintf( emptyfile, "empty %s: ", fname );
351
352     if( file == NULL || ( ch = fgetc( file ) ) == EOF ) {
353         fputs( "FILE ALREADY EMPTY\n", emptyfile );
354         if ( ch == EOF ) { fclose( file ); }
355     }
356     else {
357         // appending to empty.txt
358         do {
359             fputc( ch, emptyfile );
360         } while ( ( ch = fgetc( file ) ) != EOF );
361
362         fputc( '\n', emptyfile );
363         fclose( file );
364
365         // emptying file
366         file = fopen( fname, "w" );
367         fclose( file );
368     }
369
370     fclose( emptyfile );
371     pthread_mutex_unlock( &elock );
372
373     // sleep
374     srand(time(0));
375     int r = rand()%4 + 7;
376     sleep(r);
377
378     if( test_mode ) {
379         printf( "empty %s complete. slept for %d seconds.\n", fname, r);
380     }
381 }
```

---

---

**Code Block 20** finish\_job function

---

```
200 void finish_job( jobdata *data ) {
201     int i = data->fileidx;
202
203     // finish job
204     pthread_mutex_lock( &tlock[i] );
205     tjob *job = jobs[i];
206     if( job->workers == 0 ) {
207         free( job );
208         jobs[i] = NULL;
209         strcpy( files[i], " " );
210     }
211     else {
212         job->curr_job++;
213         pthread_cond_broadcast( &(job->done) );
214     }
215     pthread_mutex_unlock( &tlock[i] );
216
217     // retire worker
218     pthread_mutex_lock( &qlock );
219     threads[data->worker_ID] = 0;
220     free( data );
221     pthread_mutex_unlock( &qlock );
222 }
```

---

## 2.7 Master Thread Functions

The master thread is in charge of initializing the program, determining the correct data for a job, continuously receiving user requests, and dispatching worker threads.

- The `main()` function (Code Block 21) is where the master thread and the program as a whole begins. It declares the following variables:
  - `buf`: stores the user input as a string
  - `cpy`: stores a copy of user input for logging
  - `command`, `filepath`, and `string`: store the requested file operation, user-specified filepath, and string to be written (if any), respectively
  - `context`: a helper variable used by `strtok_r()` in tokenizing the string
  - `cmdidx`: the index of the command in the `cmds` array
  - `fileidx`: the index of the filepath in the `files` and `jobs` arrays
  - `job_ID`: represents a worker thread's order of operation on a file
  - `worker_ID`: represents a worker thread's location in the thread pool `tid`

To begin setting up the program, the `init()` function is called with command line arguments as parameters. Then, a `while` loop is entered where user input is continuously received via `getcmd()` and stored in `buf`. For every iteration in the `while` loop, the master thread performs the following, in order:

- The contents of `buf` are copied to `cpy`, then the last character of `buf` is changed from `'n'` to `' '` to enable tokenization.
- The array `string` is initialized to `" "`.
- Using `strtok_r()`, `buf` is tokenized based on the delimiter `" "` (stored in `space`), with `context` used to mark the location of the next token. The first token is stored in `command`. If this token is `"quit"`, the `while` loop exits and the program ends.
- The command is validated by checking whether it exists in the `cmds` array via `strindex()`. If the command is invalid, the `while` loop iterates. Otherwise, the command is logged to `commands.txt` via `log_command()`.
- The next token which is the user-specified filepath is stored in `filepath`. If the command is `"write"`, then another token which is the write data is stored in `string`.
- Using `strindex()`, the master thread checks whether the specified file is already active (i.e., already exists in `files`). If it is, its index in the `files` array is stored in `fileidx`. Otherwise, the file is instead registered via `register_job()` whose return value `fileidx` stores.
- Since every operation requires one thread, an available worker thread is sought via `register_worker()` and its index is stored in `worker_ID`.
- Now that the file has been registered in the `files` array, the corresponding `struct TJob` is prepared via `init_worker()`. The `count` member of this struct is stored in `job_ID` to function as the worker thread's job ID.
- The data obtained by the previous steps are stored in `data` as a `struct Jobdata` via `init_worker()`.
- If testing is enabled, then `test_print()` is called.

- Lastly, the worker thread is dispatched via `pthread_create()`, with `data` passed as a parameter. In order to let the OS know that it can immediately reclaim memory once a worker thread finishes an operation, the thread is detached.

Since the user may decide to quit while worker threads are not yet finished, the program exits with a call to `pthread_exit()`. This allows all worker threads to finish operation before terminating the program, preventing data loss.

- The function `init()` (Code Block 22) takes in the command line arguments as parameters. If `"test"` was used as an argument to the program, `test_mode` is set to `1`, enabling testing.

All locks in `tlock` are initialized, as well as the global locks `qlock`, `rlock`, and `elock`. Then, all entries in `files` are initialized to `" "` (rather than `NULL`). This is necessary since `files` is passed to `strindex()` in `main()`, and `strindex()` calls `strcmp()` which has undefined behavior when dealing with `NULL` strings. Lastly, the log files `commands.txt`, `read.txt`, and `empty.txt`, are emptied via `fopen()` (then `fclose()`) calls.

- The function `register_job()` (Code Block 25) seeks a free slot in `files` given a filepath by finding the first entry stored as `" "`. If a slot is found, the filepath is stored in `files` and its index is returned. Otherwise, the `files` array is full (i.e., there are currently `N` active files). In this case, the program terminates via `pthread_exit()`.
- The function `init_job()` (Code Block 24) takes in the index `idx` of a file in `files` as well as its filepath as parameters. Since a worker thread may call `finish_job()` at the same time that the master thread calls `init_job()` and since both functions modify `files` and `jobs`, `tlock[idx]` is acquired to prevent race conditions.

Note that at the start of the function (line 275), the filepath is stored again in `files` even though an earlier call from `main()` to `register_job()` already does this. This is necessary since `finish_job()` clears `files[idx]` when all worker threads are done. In between `register_job()` and `init_job()` calls is a nonzero amount of time wherein a lone worker thread may possibly call `finish_job()`. This would result in an error in the program where an active file is not stored in `files` but its `struct TJob` exists. This would cause the next worker thread to try to modify a file named `" "`, which is not intended. To prevent this very error, `strcpy()` is again called to redundantly save the file in `files`.

Afterwards, the function simply checks whether `jobs[idx]` already has a `struct TJob`. if it does, its `count` member is incremented. otherwise, memory is allocated for a `struct TJob`. first, its locks are initialized. the `workers` member is initialized to `0` since no worker threads have enqueued themselves yet. since the job id of the first worker thread would be `1`, both `count` and `curr_job` are initialized to `1`. this is necessary to prevent a deadlock in `dispatch_worker()` where the first worker never wakes up. Lastly, the struct is stored in `jobs`.

Afterwards, `tlock[idx]` is released, and `count` is returned as the worker thread's job ID.

- The function `register_worker()` (Code Block 25) seeks an available thread in `tid` by finding the first entry in `threads` equal to `0`. If a thread is found, its index `i` is returned and `threads[i]` is set to `1`. Since a thread may call `finish_job()` at the same time that the master thread calls `register_worker()` and since both functions modify `threads`, `qlock` is acquired to prevent race conditions.
- The function `init_worker()` (Code Block 26) takes in the data from `main()` equivalent to all of a `struct Jobdata`'s parameters. The function simply allocates memory for the struct, initializes each of its members with its corresponding parameter, and returns the struct's pointer.

---

**Code Block 21** main function

---

```
585 int main( int argc, char *argv[] ) {
586     char buf[MAXLEN]; // buffer where user input is stored
587     char cpy[MAXLEN]; // copy of user input
588     char command[10], filepath[51], string[51];
589     char *context;    // used by strtok to save state
590     int cmdidx;        // index of command in cmds array
591     int fileidx;       // index of file in files/jobs array
592     int job_ID;        // position of worker in job queue
593     int worker_ID;     // position of worker in thread pool
594
595     init( argc, argv );
596
597     while( getcmd( buf, sizeof( buf ) ) == 0 ) {
598         strcpy( cpy, buf );           // reset string
599         buf[strlen( buf ) - 1] = ' '; // turn '\n' to space
600         strcpy( string, " " );       // reset string
601
602         strcpy( command, strtok_r( buf, space, &context ));
603
604         if( strcmp( command, "quit" ) == 0 ) { break; }
605
606         // validating command
607         if(( cmdidx = strindex( command, cmds, 3 )) == -1 ) {
608             printf( "Invalid command.\n" ); continue;
609         }
610         log_command( cpy );
611
612         // save filepath and/or string
613         strcpy( filepath, strtok_r( NULL, space, &context ));
614         if( cmdidx == 1 ) {
615             strcpy( string, strtok_r( NULL, space, &context ));
616         }
617
618         // check if job already exists, else save it
619         if(( fileidx = strindex( filepath, files, N )) == -1 ) {
620             fileidx = register_job( filepath );
621         }
622
623         // find an available thread in thread pool
624         worker_ID = register_worker();
625
626         // ready jobs and worker data
627         job_ID = init_job( fileidx, filepath );
628         jobdata *data = init_worker( cmdidx, fileidx, job_ID, worker_ID, string );
629
630         // print job and worker status for testing
631         if( test_mode ) { test_print( fileidx, data ); }
632
633         // dispatch worker
634         pthread_create( &tid[worker_ID], NULL, dispatch_worker, ( void * )data );
635         pthread_detach( tid[worker_ID] );
636     }
637
638     pthread_exit(NULL);
639 }
```

---



---

**Code Block 22** init function

---

```
439 void init( int argc, char *argv[] ) {
440
441     if( argc == 2 ) {
442         test_mode = strcmp( argv[1], "test" ) ? 0 : 1;
443     }
444
445     for( int i = 0; i < N; i++ ) {
446         if( pthread_mutex_init( &tlock[i], NULL ) != 0 ) {
447             printf( "mutex init has failed\n" );
448         }
449     }
450
451     if( pthread_mutex_init( &qlock, NULL ) != 0 ||
452         pthread_mutex_init( &rlock, NULL ) != 0 ||
453         pthread_mutex_init( &elock, NULL ) != 0 ) {
454         printf( "mutex init has failed\n" );
455     }
456
457     for( int i = 0; i < N; i++ ) {
458         strcpy( files[i], " " );
459     }
460
461     cmdfile = fopen( "commands.txt", "w" );
462     readfile = fopen( "read.txt", "w" );
463     emptyfile = fopen( "empty.txt", "w" );
464     fclose( cmdfile );
465     fclose( readfile );
466     fclose( emptyfile );
467
468     printf( "Initialization complete. Type \"quit\" to exit the program.\n" );
469 }
```

---

---

**Code Block 23** register\_job function

---

```
480 int register_job( char *filepath ) {
481     for( int i = 0; i < N; i++ ) {
482         if( strcmp( files[i], " " ) == 0 ) {
483             strcpy( files[i], filepath );
484             return i;
485         }
486     }
487     // no available slots
488     printf( "Too many jobs at once. Program will exit after all operations "
489         "are finished.\n" );
490     pthread_exit(NULL);
491 }
```

---

---

**Code Block 24** init\_job function

---

```
505 int init_job( int idx, char *filepath ) {
506     pthread_mutex_lock( &tlock[idx] );
507
508     strcpy( files[idx], filepath );
509     tjob* job = jobs[idx];
510
511     // update job
512     if( job != NULL ) {
513         job->count++;
514     }
515     // create job
516     else {
517         job = malloc( sizeof( tjob ) );
518
519         if( pthread_mutex_init( &( job->wlock ), NULL ) != 0 ||
520             pthread_mutex_init( &( job->global_lock ), NULL ) != 0 ) {
521             printf( "mutex init has failed\n" );
522         }
523
524         if(pthread_cond_init( &( job->done ), NULL ) != 0) {
525             printf( "cond init has failed" );
526         }
527
528         job->workers = 0;
529         job->count = 1;
530         job->curr_job = 1;
531         jobs[idx] = job;
532     }
533
534     pthread_mutex_unlock( &tlock[idx] );
535     return job->count;
536 }
```

---

---

**Code Block 25** register\_worker function

---

```
546 int register_worker() {
547     pthread_mutex_lock( &qlock );
548     for( int i = 0; i < M; i++ ) {
549         if( threads[i] == 0 ) {
550             threads[i] = 1;
551             pthread_mutex_unlock( &qlock );
552             return i;
553         }
554     }
555     pthread_mutex_unlock( &qlock );
556     // no available threads
557     printf( "No more available threads. Program will exit after all "
558         "operations are finished.\n" );
559     pthread_exit(NULL);
560 }
```

---

---

**Code Block 26** init\_worker function

---

```
575 jobdata *init_worker( int cmdidx, int fileidx, int j_ID, int w_ID, char *str ) {  
576     jobdata *data = malloc( sizeof(jobdata) );  
577     data->cmdidx = cmdidx;  
578     data->fileidx = fileidx;  
579     data->j_ID = j_ID;  
580     data->worker_ID = w_ID;  
581     strcpy( data->string, str );  
582     return data;  
583 }
```

---

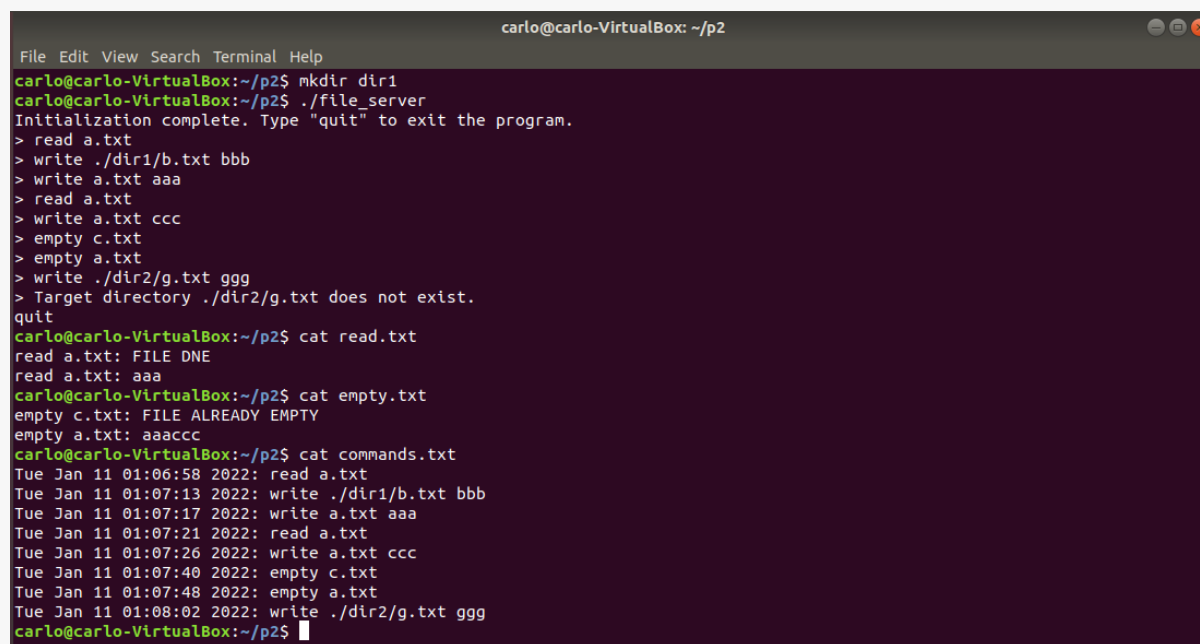
## 3 Testing the System

### 3.1 Correctness of Commands

Figure 1 shows a run of the program with testing disabled. The subdirectory `dir1` is created beforehand, but no files exist yet. Notice that the command `write ./dir2/g.txt ggg` is not carried out. This is because the subdirectory does not exist, and the program is not designed to handle this situation.

As shown in the log files, we can see that all commands work as intended.

Figure 1. Sample program run



```
carlo@carlo-VirtualBox: ~/p2
File Edit View Search Terminal Help
carlo@carlo-VirtualBox:~/p2$ mkdir dir1
carlo@carlo-VirtualBox:~/p2$ ./file_server
Initialization complete. Type "quit" to exit the program.
> read a.txt
> write ./dir1/b.txt bbb
> write a.txt aaa
> read a.txt
> write a.txt ccc
> empty c.txt
> empty a.txt
> write ./dir2/g.txt ggg
> Target directory ./dir2/g.txt does not exist.
quit
carlo@carlo-VirtualBox:~/p2$ cat read.txt
read a.txt: FILE DNE
read a.txt: aaa
carlo@carlo-VirtualBox:~/p2$ cat empty.txt
empty c.txt: FILE ALREADY EMPTY
empty a.txt: aaaccc
carlo@carlo-VirtualBox:~/p2$ cat commands.txt
Tue Jan 11 01:06:58 2022: read a.txt
Tue Jan 11 01:07:13 2022: write ./dir1/b.txt bbb
Tue Jan 11 01:07:17 2022: write a.txt aaa
Tue Jan 11 01:07:21 2022: read a.txt
Tue Jan 11 01:07:26 2022: write a.txt ccc
Tue Jan 11 01:07:40 2022: empty c.txt
Tue Jan 11 01:07:48 2022: empty a.txt
Tue Jan 11 01:08:02 2022: write ./dir2/g.txt ggg
carlo@carlo-VirtualBox:~/p2$
```

### 3.2 Sequential Execution of Requests to the Same File

Figure 2 shows a run of the program with testing enabled. Recall that when testing is enabled, the following information is printed per request:

- All active files and all active worker threads
- The specified file's `struct TJob` members
- The worker thread's `struct Jobdata` members

The first command is "`read a.txt`" (1), which creates a new entry in `files`. Note that this file does not exist yet. Thus, a `struct TJob` is stored in `jobs`, with the expected initial values of its members. Since the worker thread's job ID is the same as the job's `curr_job`, it immediately begins operation.

As this is happening, a second command "`write a.txt fasjd`" is entered (2). The `count` member of the job is updated, and the job ID of the new worker thread is 2.

A third "read" command to `a.txt` is supplied (3). Meanwhile, we find that the thread with a `worker_ID` of 1 (the second worker thread) is done sleeping (4). Since the first worker thread is not yet done reading, this thread sleeps again. Notice that in (5), the thread with a `worker_ID` of 0 has finished sleeping and reading. Only then does the second thread continue execution.

Lastly, an “empty” command (6) is entered. Notice that it is not carried out until the preceding read is finished (7) since, again, the worker thread’s `job_ID` (which is 4) and the job’s `curr_job` (which is 3) are not equal at the time that the command is issued.

Figure 3 shows that the output of the log files are as expected. Notice that the first read to `a.txt` logs `FILE DNE`. After the write, the log is now `"fasjd"`. Since the last operation on `a.txt` was empty, running `cat a.txt` produces nothing.

With this test, we have shown the following:

- Requests to the same file are carried out sequentially despite worker threads exiting `xsleep()` at unpredictable times, which prevents simultaneous operations to the same file.
- The master thread can continuously receive requests without blocking.

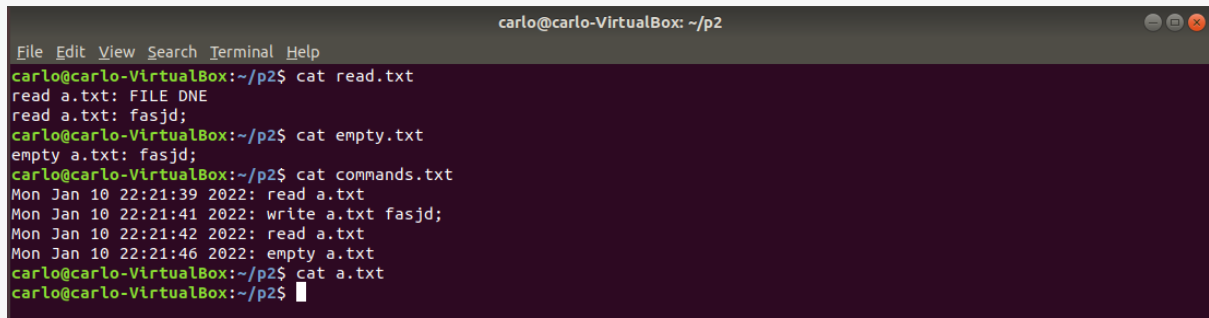
Figure 2. Submitting multiple requests to one file

```

carlo@carlo-VirtualBox: ~/p2
File Edit View Search Terminal Help
carlo@carlo-VirtualBox:~/p2$ ./file_server test
Initialization complete. Type "quit" to exit the program.
> read a.txt
files[0]: a.txt
active threads: 0,
file <a.txt> currently has:
  workers = 0
  count = 1
  curr_job = 1
thread working on fileidx = 0 with:
  command = read
  job_ID = 1
  worker_ID = 0
  string =
----- (1)
> write a.txt fasjd;
files[0]: a.txt
active threads: 0, 1,
file <a.txt> currently has:
  workers = 1
  count = 2
  curr_job = 1
thread working on fileidx = 0 with:
  command = write
  job_ID = 2
  worker_ID = 1
  string = fasjd;
----- (2)
> read a.txt
thread 1 slept for 1 seconds.
files[0]: a.txt
active threads: 0, 1, 2,
file <a.txt> currently has:
  workers = 2
  count = 3
  curr_job = 1
thread working on fileidx = 0 with:
  command = read
  job_ID = 3
  worker_ID = 2
  string =
----- (3)
----- (4)
> thread 0 slept for 6 seconds.
read a.txt complete.
write a.txt complete. slept for 150 milliseconds.
empty a.txt
files[0]: a.txt
active threads: 0, 2,
file <a.txt> currently has:
  workers = 1
  count = 4
  curr_job = 3
thread working on fileidx = 0 with:
  command = empty
  job_ID = 4
  worker_ID = 0
  string =
----- (5)
----- (6)
> thread 2 slept for 6 seconds.
read a.txt complete.
thread 0 slept for 6 seconds.
empty a.txt complete. slept for 9 seconds.
quit
carlo@carlo-VirtualBox:~/p2$

```

Figure 3. Execution results



```
carlo@carlo-VirtualBox: ~/p2
File Edit View Search Terminal Help
carlo@carlo-VirtualBox:~/p2$ cat read.txt
read a.txt: FILE DNE
read a.txt: fasjd;
carlo@carlo-VirtualBox:~/p2$ cat empty.txt
empty a.txt: fasjd;
carlo@carlo-VirtualBox:~/p2$ cat commands.txt
Mon Jan 10 22:21:39 2022: read a.txt
Mon Jan 10 22:21:41 2022: write a.txt fasjd;
Mon Jan 10 22:21:42 2022: read a.txt
Mon Jan 10 22:21:46 2022: empty a.txt
carlo@carlo-VirtualBox:~/p2$ cat a.txt
carlo@carlo-VirtualBox:~/p2$
```

### 3.3 Concurrent Execution of Requests to Different Files

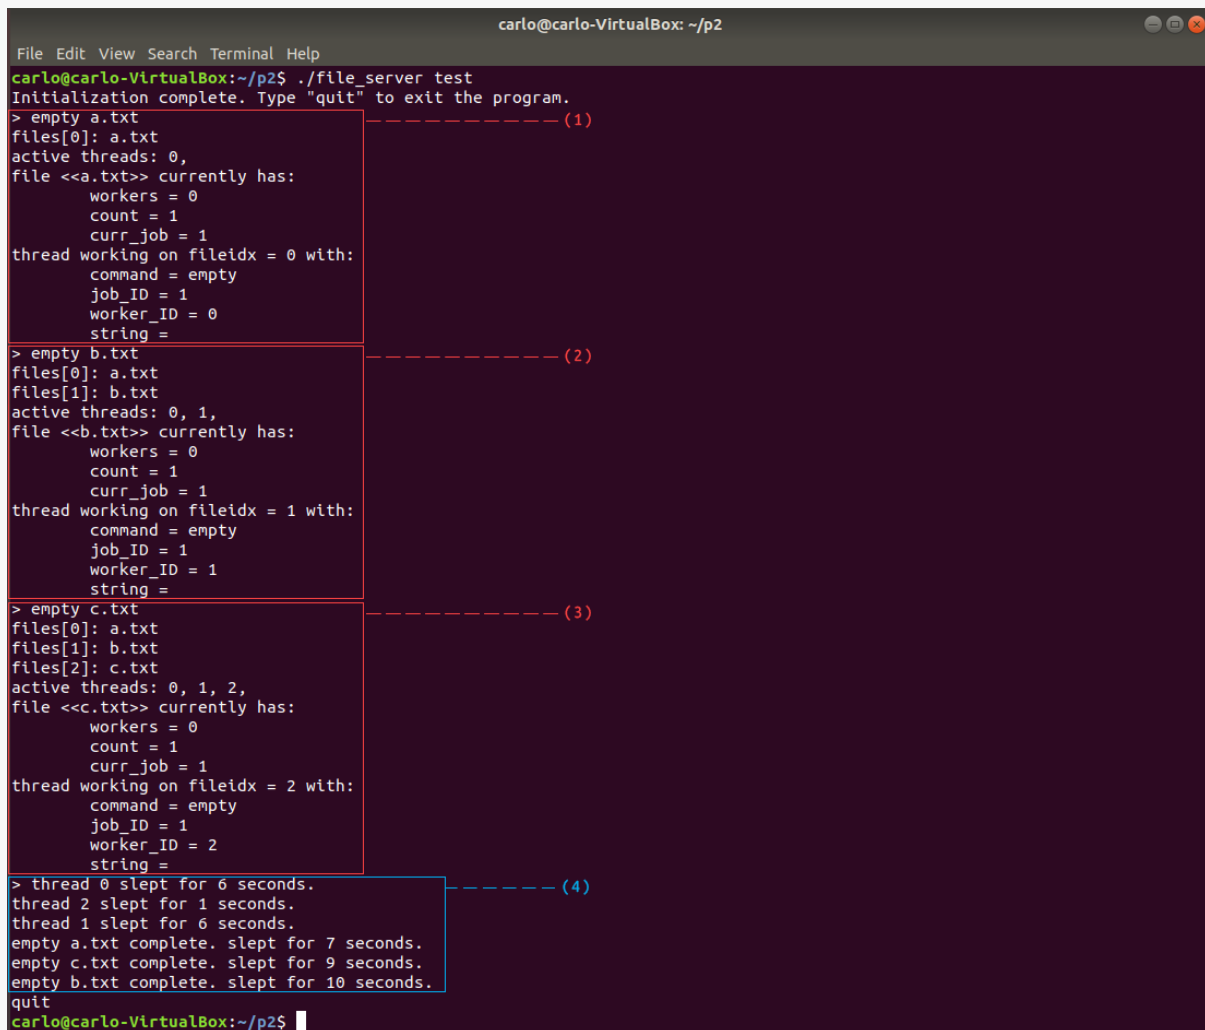
Figure 4 shows a run of the program with testing enabled. Three “empty” commands are issued (1–3). Note that for each of these commands, the list of active files is printed, and that each command registers a new file. This allows the program to perform these operations concurrently.

From (4) and Figure 5, we can see that the following are all different:

- The order in which the commands were issued
- The order in which the worker threads finished sleeping
- The order in which the operations were logged to `empty.txt`

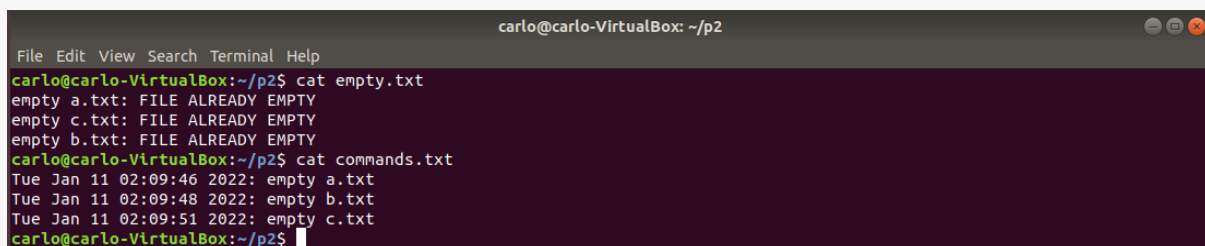
This shows that since no order is enforced, the end result of concurrent operations is non-deterministic.

Figure 4. Submitting multiple requests to different files



```
carlo@carlo-VirtualBox: ~/p2
File Edit View Search Terminal Help
carlo@carlo-VirtualBox:~/p2$ ./file_server test
Initialization complete. Type "quit" to exit the program.
> empty a.txt ----- (1)
files[0]: a.txt
active threads: 0,
file <<a.txt>> currently has:
  workers = 0
  count = 1
  curr_job = 1
thread working on fileidx = 0 with:
  command = empty
  job_ID = 1
  worker_ID = 0
  string =
> empty b.txt ----- (2)
files[0]: a.txt
files[1]: b.txt
active threads: 0, 1,
file <<b.txt>> currently has:
  workers = 0
  count = 1
  curr_job = 1
thread working on fileidx = 1 with:
  command = empty
  job_ID = 1
  worker_ID = 1
  string =
> empty c.txt ----- (3)
files[0]: a.txt
files[1]: b.txt
files[2]: c.txt
active threads: 0, 1, 2,
file <<c.txt>> currently has:
  workers = 0
  count = 1
  curr_job = 1
thread working on fileidx = 2 with:
  command = empty
  job_ID = 1
  worker_ID = 2
  string =
> thread 0 slept for 6 seconds. ----- (4)
thread 2 slept for 1 seconds.
thread 1 slept for 6 seconds.
empty a.txt complete. slept for 7 seconds.
empty c.txt complete. slept for 9 seconds.
empty b.txt complete. slept for 10 seconds.
quit
carlo@carlo-VirtualBox:~/p2$
```

Figure 5. Execution results



```
carlo@carlo-VirtualBox: ~/p2
File Edit View Search Terminal Help
carlo@carlo-VirtualBox:~/p2$ cat empty.txt
empty a.txt: FILE ALREADY EMPTY
empty c.txt: FILE ALREADY EMPTY
empty b.txt: FILE ALREADY EMPTY
carlo@carlo-VirtualBox:~/p2$ cat commands.txt
Tue Jan 11 02:09:46 2022: empty a.txt
Tue Jan 11 02:09:48 2022: empty b.txt
Tue Jan 11 02:09:51 2022: empty c.txt
carlo@carlo-VirtualBox:~/p2$
```